

Solution: (a) We use a layered construction to keep track of whether Yulie has bought empanadas and whether a gas station has been reached. For the layer where a gas station has not been reached, we need to know which vertices can be reached and which cannot. Therefore, we need to run Dijkstra's algorithm once on G not taking gas stations into account. Following is the algorithm.

- (1) Run Dijkstra's algorithm on G starting at s to compute $pred(v)$, the predecessor of v in the shortest path from s to v . Let $remainingMiles(v)$ be the maximum number of miles that Yulie's car can run when she reaches v . It can be defined as follows:

$$remainingMiles(v) = \begin{cases} D & v = s \\ remainingMiles(p) - l(pv) & \text{otherwise, } p = pred(v) \end{cases}$$

The time complexity of computing $remainingMiles$ is $O(1)$ and we can put its computation into each iteration of Dijkstra's algorithm. Since Dijkstra's algorithm explores the shortest path, $remainingMiles(v)$ must be the maximum remaining number of miles at v .

- (2) Construct a graph $G' = (V', E')$, and define edge lengths $l'(e)$ for E' :
- $V' = V \times \{NE, E\} \times \{NG, G\}$, where NE means Yulie has not bought empanadas yet, E means Yulie has bought empanadas, NG means a gas station has not been reached, G means a gas station has been reached.
 - For each edge $uv \in E$, we add the following to E' and $l'(e)$. We use a and b to generalize empanadas and gas stations.
 - $(u, a, NG)(v, a, NG)$ if $remainingMiles(v) \geq 0$ and $u = pred(v)$; $l'(e) = l(uv)$
 - $(u, a, G)(v, a, G)$; $l'(e) = l(uv)$
 - $(u, a, NG)(u, a, G)$ if $u \in Y$; $l'(e) = 0$
 - $(u, NE, b)(u, E, b)$ if $u \in X$; $l'(e) = 0$
- (3) Let $s' = (s, NE, NG)$. Run Dijkstra's algorithm on G' starting at s' using $l'(e)$ as edge weights. Let $dist(v)$ be the shortest distance from s' to v in G' . Return the minimum of $dist((t, E, NG))$ and $dist((t, E, G))$.

It is easy to also return a shortest path since Dijkstra's algorithm can keep track of the path. Note that $dist(v) = \infty$ if there is not a path from s to v in Dijkstra's algorithm. Therefore the algorithm will return ∞ if there is no way to reach the destination. In the new graph G' , $|V'| = 4|V|$ and $|E'| \leq |X| + |Y| + 2|E| \leq |V| + 2|E|$. Since the running time for Dijkstra's algorithm is $O(|E| + |V| \log |V|)$, the running time of this algorithm is $O(|V| + 2|E| + 4|V| \log 4|V|) = O(m + n \log n)$.

- (b) Let $k = |Y|$. We want to know the "reach" of each gas station to account for the fact that Yulie can only run R miles after reaching a gas station. Therefore, we run Dijkstra's algorithm k times using each gas station as the start node. The layers are concatenated and form a new graph. Then the problem is reduced to an SSSP problem in the new graph. The algorithm is as follows:

- (1) Run Dijkstra's algorithm $k + 1$ times on G starting from s, y_1, y_2, \dots, y_k where $y_i \in Y$. Let $\text{pred}(y_i, v)$ be the predecessor of v in the shortest path from y_i to v . Each time Dijkstra's algorithm is run starting from y_i , we compute an array $\text{remainingMiles}(y_i, v)$ using the same technique as in (a) which represents the remaining miles when Yulie reaches v from y_i . It is defined as follows:

$$\text{remainingMiles}(y_i, v) = \begin{cases} R & v = y_i \\ \text{remainingMiles}(y_i, p) - l(pv) & \text{otherwise, } p = \text{pred}(y_i, v) \end{cases}$$

Let G_1, G_2, \dots, G_k be the resulting DAGs representing the shortest paths starting from y_1, y_2, \dots, y_k . Then for each G_i , we remove all edges $\{uv \mid \text{remainingMiles}(y_i, v) < 0, u = \text{pred}(y_i, v)\}$.

- (2) Construct a new graph $G' = (V', E')$:
- $V' = V \times \{N, R_1, R_2, \dots, R_k\} \times \{E, NE\}$. E means Yulie has bought empanadas and NE means Yulie has not bought empanadas. N means Yulie has not reached any gas station yet, R_i means the last gas station reached is y_i .
 - For each edge $uv \in E$, we add the following to E' and $l'(e)$. We use r to generalize $\{N, R_1, R_2, \dots, R_k\}$ and a to generalize $\{E, NE\}$.
 - $(u, N, a)(v, N, a)$ if $\text{remainingMiles}(s, v) \geq 0$ and $u = \text{pred}(s, v)$; $l'(e) = l(uv)$
 - $(u, R_i, a)(v, R_i, a)$ if uv is an edge in G_i ; $l'(e) = l(uv)$
 - $(u, r, NE)(u, r, E)$ if $u \in X$; $l'(e) = 0$
 - $(u, r, a)(u, R_i, a)$ if $u = y_i$; $l'(e) = 0$
- (3) Let $s' = (s, N, NE)$ and run Dijkstra's algorithm on G' starting from s' using $l'(e)$ as edge weights. Return the minimum of the lengths of the shortest paths from s' to (t, r, E) where $r \in \{N, R_1, R_2, \dots, R_k\}$.

The algorithm will return ∞ if there is no valid path due to the implementation of Dijkstra's algorithm. The size of the new graph G' is $|V'| = 2(k + 1)n$ and $|E'| \leq (k + 1)m + n$. We run Dijkstra $k + 1$ times on the original graph and once on G' , so the overall time complexity is $O(km + kn \log n)$ where $k = |Y|$.

■

Solution: (a) We begin by computing the meta-graph G^{SCC} of G . The weight of each node $w'(u)$ in G^{SCC} is the sum of all $w(v)$ where $v \in u$ because the eggs at each node can only be collected once. We want to use the algorithm to find the longest path in a DAG (which is mentioned in lecture). In order to convert vertex weights to edge weights, we divide each vertex $v \in G^{\text{SCC}}$ into two halves v_{in} and v_{out} with v_{in} connected to the in-edge and v_{out} connected to the out-edge. The weights of these two edges are 0. We then add an edge $v_{\text{in}}v_{\text{out}}$ with weight $l(e) = w'(v)$. Following is the algorithm.

- (1) Compute meta-graph $G^{\text{SCC}} = (V^{\text{SCC}}, E^{\text{SCC}})$ using DFS. Let $w'(u)$ be the weight of vertex u in V^{SCC} . For each $u \in V^{\text{SCC}}$, compute $w'(u) = \sum_{v \in u} w(v)$.
- (2) Construct graph $G' = (V', E')$ defined as follows:
 - $V' = V^{\text{SCC}} \times \{\text{in}, \text{out}\}$
 - $E' = \{(u, \text{out})(v, \text{in}) \mid uv \in E^{\text{SCC}}\} \cup \{(u, \text{in})(u, \text{out}) \mid u \in V^{\text{SCC}}\}$
 - For all edges in the form $e = (u, \text{in})(u, \text{out})$, edge weight $l(e) = w'(u)$. For the other edges, $l(e) = 0$.
- (3) We can see G' is a DAG because splitting vertices cannot create cycles. Let $s' = (s, \text{in})$. Run the algorithm mentioned in lecture to find the longest path in G' starting from s' . Let $d(v)$ be the length of the longest path between s' and v . Find the largest $d(v)$ among all $v \in V'$. It is equal to the maximum number of eggs that my friends can collect starting from s .

The size of G' is $|V'| = 2|V^{\text{SCC}}|$ and $|E'| = |E^{\text{SCC}}| + |V^{\text{SCC}}|$. Since the algorithm for finding the longest path in a DAG runs in $O(m + n)$ time and computing meta-graph also runs in $O(m + n)$ time, this algorithm runs in $O(m + n)$ time where $m = |E|$ and $n = |V|$.

- (b) We begin by topologically sorting G^{SCC} . We observe that if uv is an edge in the topological sort, then the length of the longest path starting from u must be larger than the one starting from v . Therefore, for all vertices reachable from u , we only need to calculate once. Following is the algorithm.

MaxEggs(G):

```

compute meta-graph  $G^{\text{SCC}}$ 
do DFS to get a topological sort of  $G^{\text{SCC}}$ 
mark all vertices in  $G^{\text{SCC}}$  as unvisited
 $\text{maxEggs} \leftarrow 0$ 
for each  $v \in G^{\text{SCC}}$  in topological order
  if  $v$  is unvisited
    run the algorithm described in (a) using  $v$  as the start vertex
     $\text{maxEggs} \leftarrow \max(\text{maxEggs}, \text{computed maximum number of eggs starting at } v)$ 
do a whatever-first search on  $G^{\text{SCC}}$ , mark all vertices reachable from  $v$  as visited

```

We can see that although we potentially run the algorithm in (a) multiple times, the parameter of each run is a smaller subgraph of G^{SCC} , and the sum of the sizes of the subgraphs is exactly equal to the size of G^{SCC} . Since the algorithm in (a) runs in $O(m + n)$ time, the algorithm above runs in $O(m + n) + \sum_i O(m_i + n_i)$ time where $\sum_i m_i = m$ and $\sum_i n_i = n$. Therefore, the overall time complexity is still $O(m + n)$.

(c) For this problem, we use G^{SCC} and vertex weights $w'(u) = \sum_{v \in u} w(v)$ for vertex $u \in G^{\text{SCC}}$. Let s be an arbitrary start node. Let $\text{maxEggs}(v, k)$ be the maximum possible number of eggs collected at at most k vertices in the path from s to v . Consider the following three cases:

- We have collected eggs at less than k locations. Then $\text{maxEggs}(v, k) = \text{maxEggs}(v, k - 1)$.
- We have collected eggs at exactly k locations, and we collect eggs at v . Then $\text{maxEggs}(v, k) = \max_{uv \in E^{\text{SCC}}} (\text{maxEggs}(u, k - 1) + w'(v))$.
- We have collected eggs at exactly k locations, and we do not collect eggs at v . Then $\text{maxEggs}(v, k) = \max_{uv \in E^{\text{SCC}}} (\text{maxEggs}(u, k))$.

The base case is when $k = 0$ and we cannot get any eggs. Also, $\text{maxEggs}(s, k) = w'(s)$ for all $k > 0$ because we can only collect eggs at s . Following are the recurrences.

$$\text{maxEggs}(v, k) = \begin{cases} 0 & k = 0 \\ w'(s) & v = s, k > 0 \\ \max \begin{cases} \text{maxEggs}(v, k - 1) \\ \max_{uv \in E^{\text{SCC}}} (\text{maxEggs}(u, k - 1) + w'(v)) \\ \max_{uv \in E^{\text{SCC}}} (\text{maxEggs}(u, k)) \end{cases} & \text{otherwise} \end{cases}$$

The idea is similar to (b): if there is an edge uv in the topological sort, then the maximum number of eggs starting from u must be larger than or equal to the maximum number of eggs starting from v , since one can always ignore the eggs in u and start from v . Therefore we only need to consider cases where we start at a vertex with no in-edge.

```

MAXEGGSATMOSTKLOCATIONS( $G, k$ ):
  compute meta-graph  $G^{\text{SCC}} = (V^{\text{SCC}}, E^{\text{SCC}})$ 
   $n \leftarrow |V^{\text{SCC}}|$ 
  integer  $\text{maxEggs}[n][k + 1]$ 
  initialize  $\text{maxEggs}[v][0] = 0$  for  $v \in V^{\text{SCC}}$ 
  for  $i \leftarrow 1$  to  $k$ 
    for  $v \in V^{\text{SCC}}$  in topological order
      if  $v$  has no in-edge
         $\text{maxEggs}[v][i] \leftarrow w'(v)$ 
      else
         $\text{maxEggs}[v][i] \leftarrow \text{maxEggs}[v][i - 1]$ 
        for each edge  $uv$  in  $E^{\text{SCC}}$ 
           $\text{maxEggs}[v][i] \leftarrow \max(\text{maxEggs}[v][i], \text{maxEggs}[u][i - 1] + w'(v))$ 
           $\text{maxEggs}[v][i] \leftarrow \max(\text{maxEggs}[v][i], \text{maxEggs}[u][i])$ 
  return  $\max_{v \in V^{\text{SCC}}} (\text{maxEggs}[v][k])$ 

```

To check if a vertex has any in-edge, we can create an $O(n)$ space array initialized with 0. Then we scan all edges uv and set the v^{th} element to 1 in $O(m)$ time. Then checking for in-edge takes $O(1)$ time for each vertex. Since we scan all vertices and all edges k times with constant-time operations during each scan, the time complexity of this algorithm is $O(k(m + n))$. ■